Biological Data CSV

Volker Hoffmann (volker@cheleb.net)

Data 1/2

Drop One?

PCA?

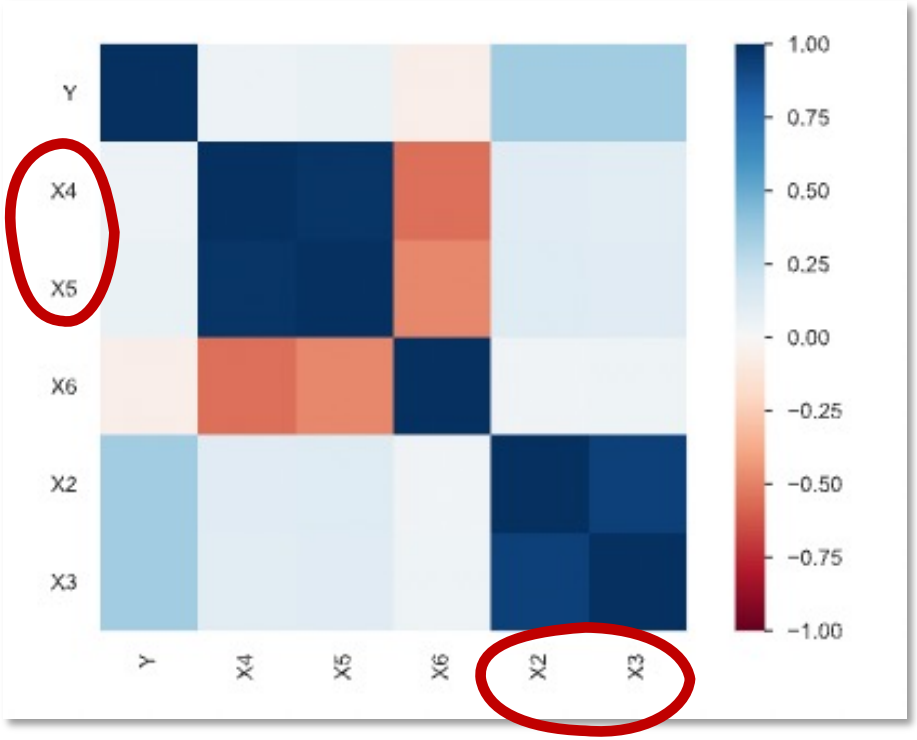

Drop One?

Drop

Drop

Alerts

X1 has constant value "human"	Constant
Dataset has 34 (1.7%) duplicate rows	Duplicates
X4 is highly overall correlated with X5 and 1 other fields	High correlation
X5 is highly overall correlated with X4	High correlation
X6 is highly overall correlated with X4	High correlation
X2 is highly overall correlated with X3	High correlation
X3 is highly overall correlated with X2	High correlation

Data 2/2

X2
Categorical

Distinct	10
Distinct (%)	0.5%
Missing	0
Missing (%)	0.0%
Memory size	31.2 KiB

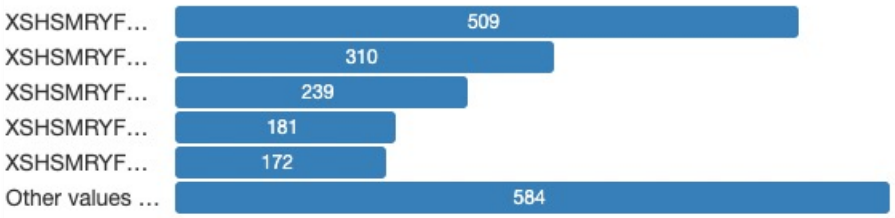

Encode as **One-Hot** / ~~Label~~
(Same for X3)

"Dummy Variable Trap"!

X5
Real number (ℝ)

Distinct	374
Distinct (%)	18.7%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	111.00359
Minimum	78.333333
Maximum	142.77778
Zeros	0
Zeros (%)	0.0%
Negative	0
Negative (%)	0.0%
Memory size	31.2 KiB

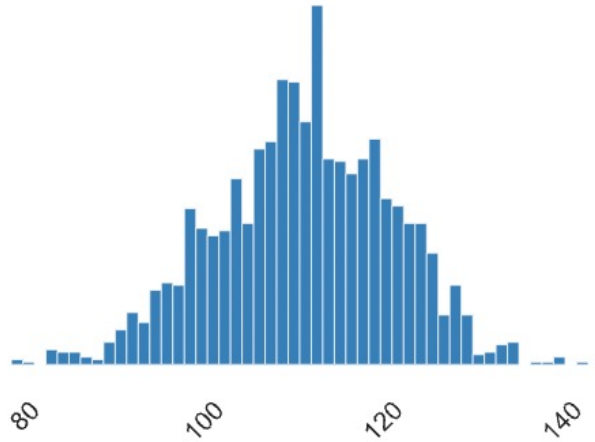

🦄 Np.rand.randn * 20 + 110 → ?
(X4, X6 also look like this)

Regression 1/3

Y

Real number (ℝ)

Distinct	557
Distinct (%)	27.9%
Missing	0
Missing (%)	0.0%
Infinite	0
Infinite (%)	0.0%
Mean	5853.5223
Minimum	0
Maximum	20000
Zeros	14
Zeros (%)	0.7%
Negative	0
Negative (%)	0.0%
Memory size	31.2 KiB

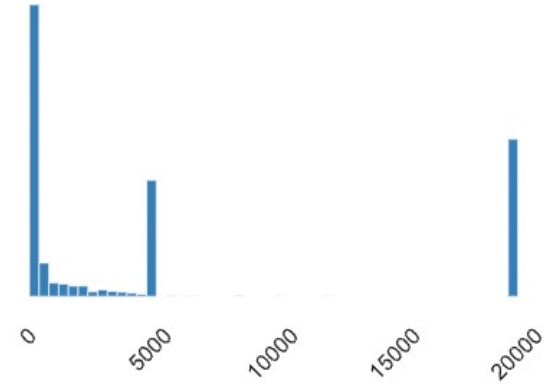

Regression 1/3

- Probably tricky… (look at histogram!)

- 3 Models
  - XGBoost
  - SVM
  - Ridge Regression (LinReg + Regularization)

- Rescaled Features to ~[0,1] (SVR/LinReg Care)
- 5-Fold Cross-Validation
- (Negative) Mean Squared Error

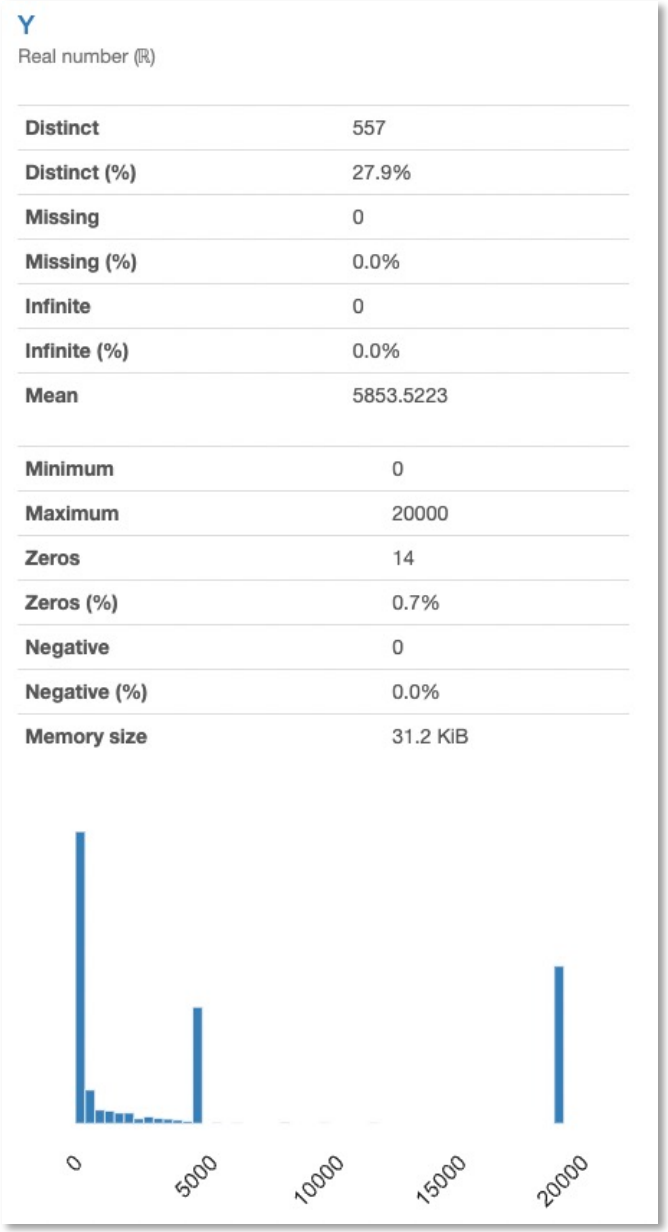

Regression 2/3

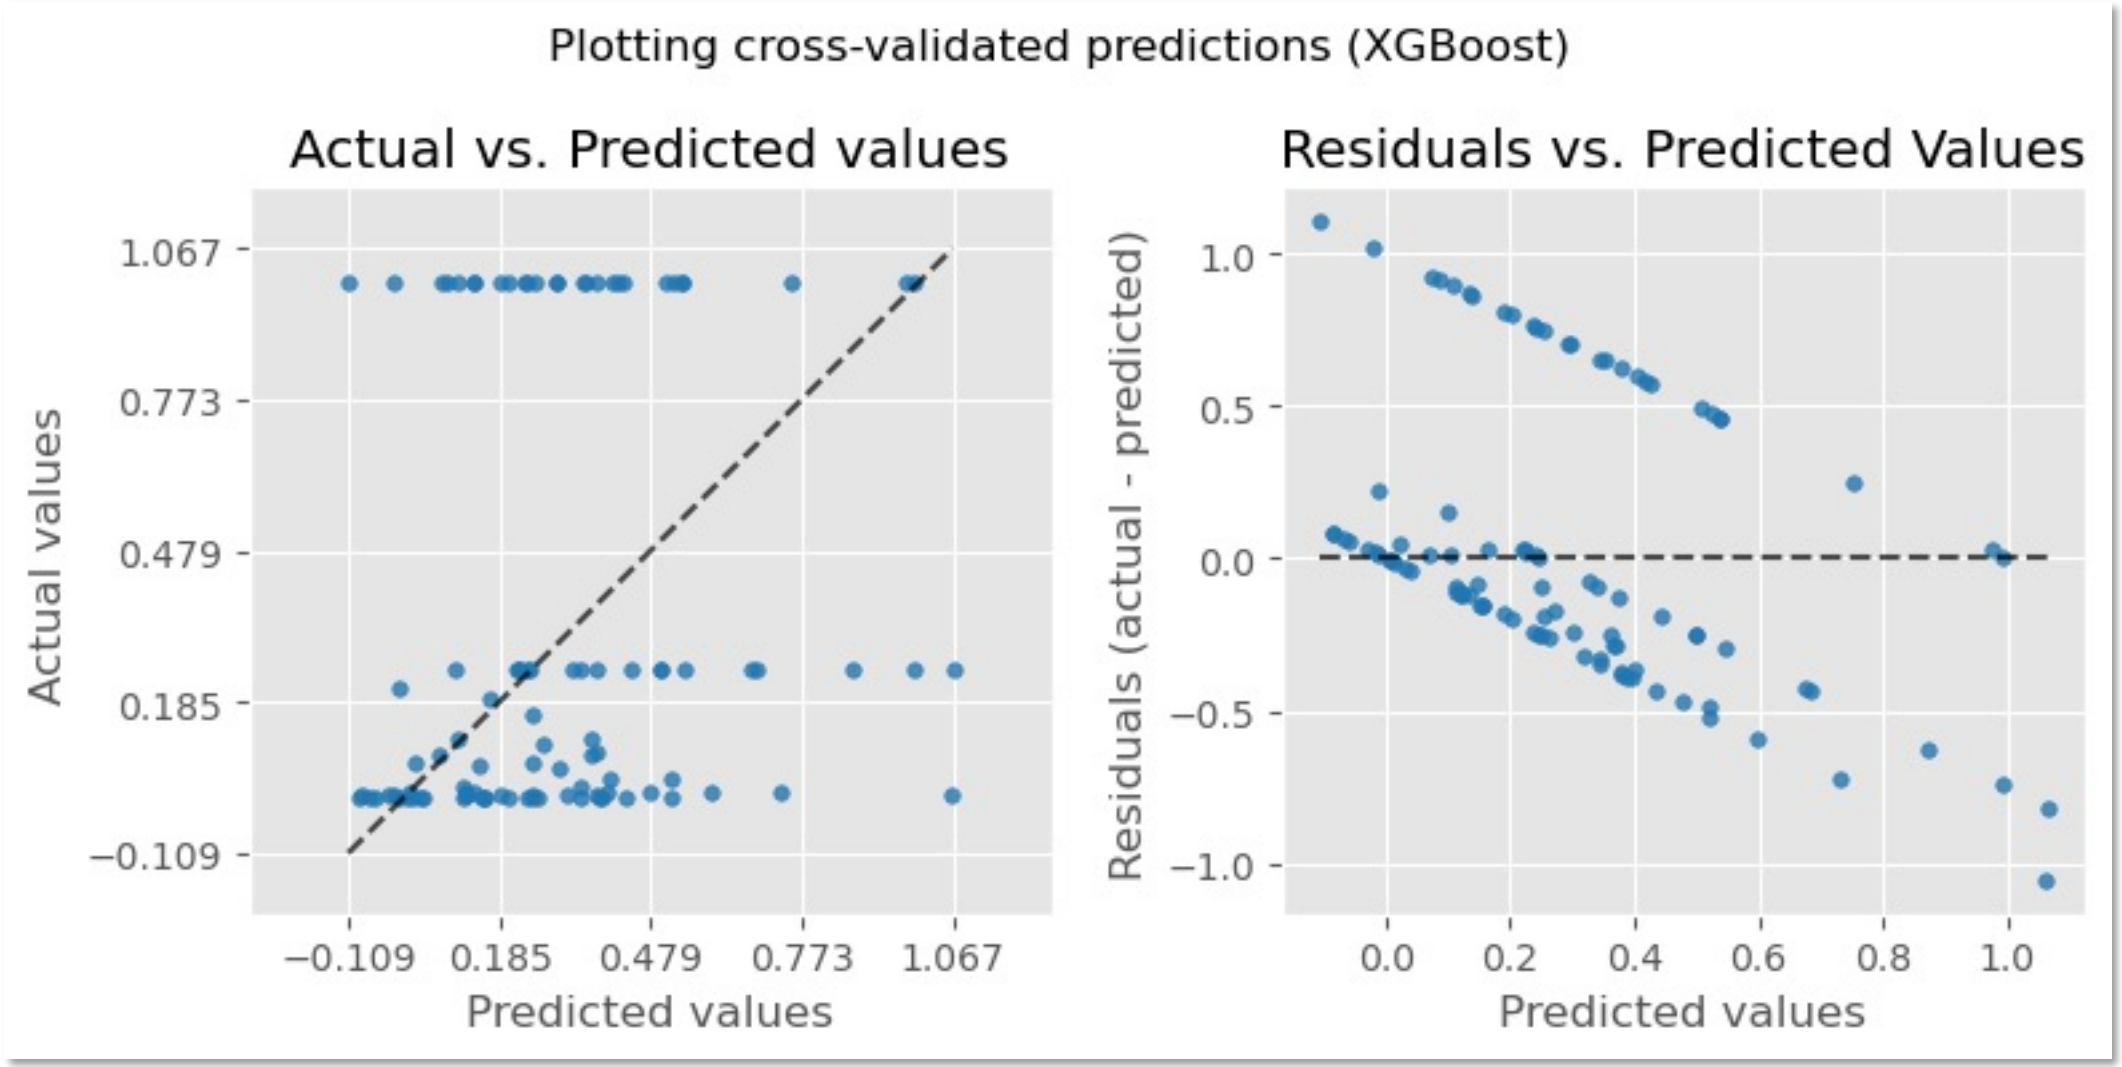

Regression 3/3

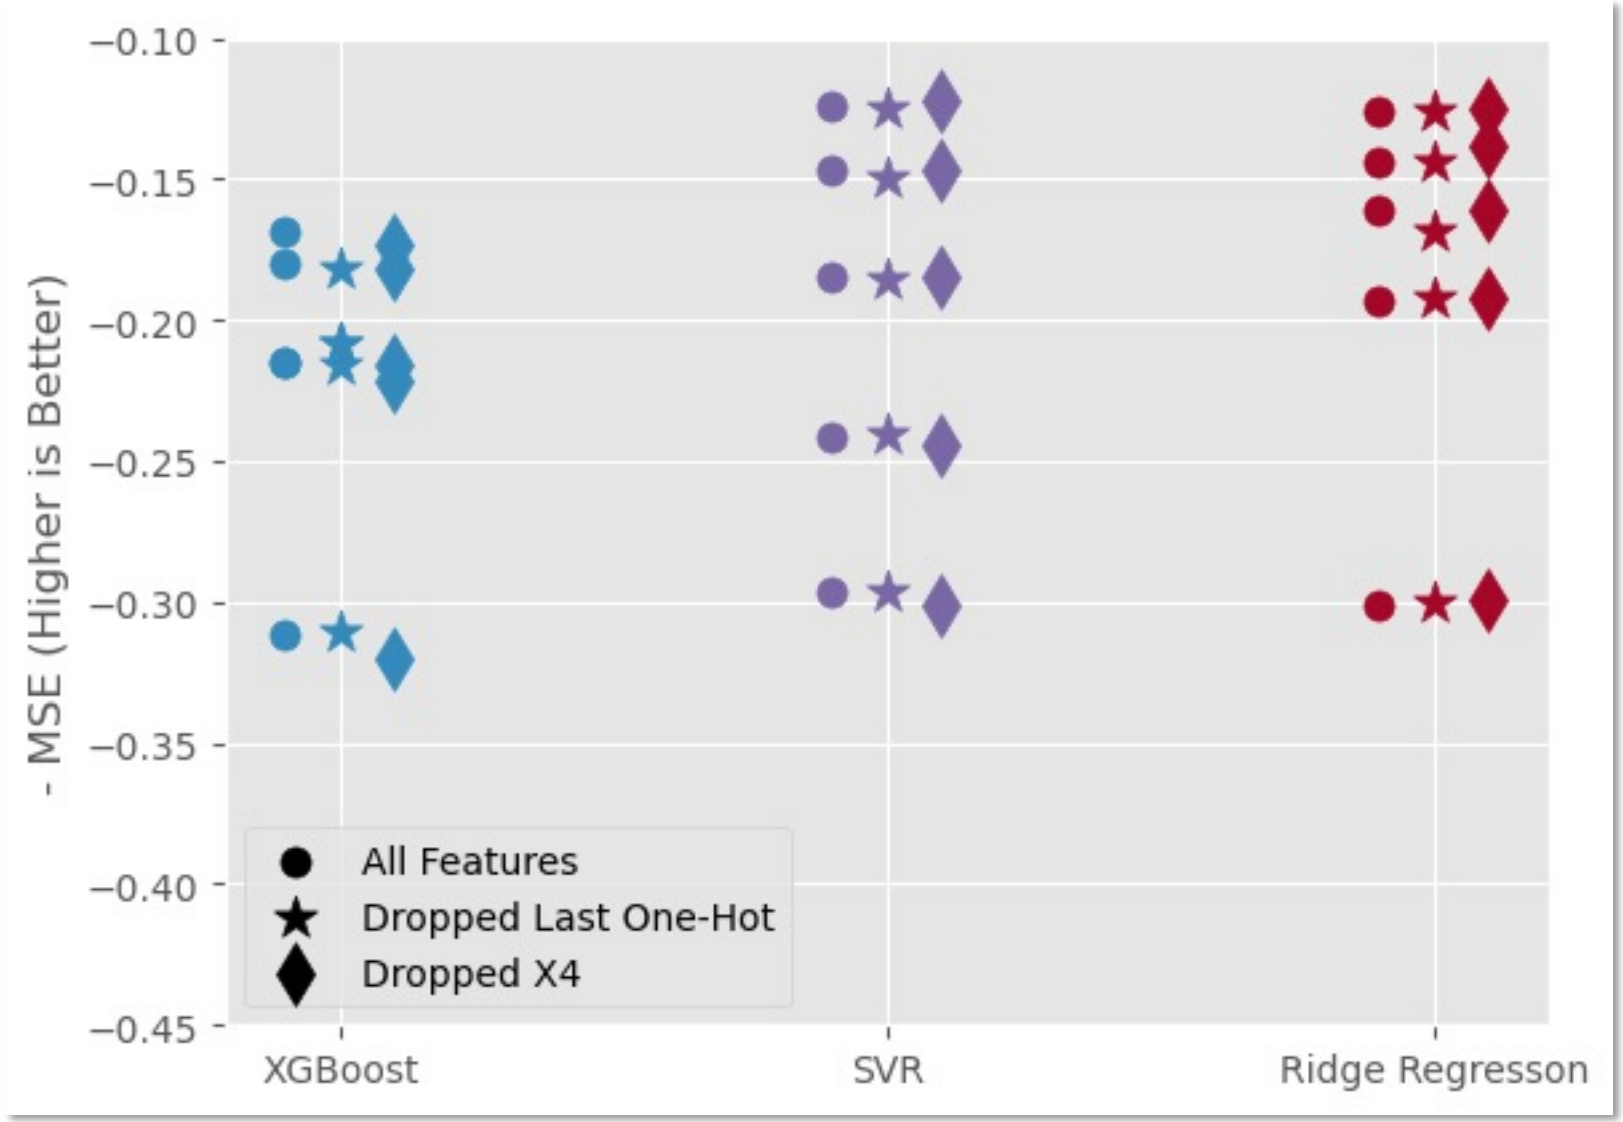

Classification 1/3

Yc

Boolean

Distinct	2
Distinct (%)	0.1%
Missing	0
Missing (%)	0.0%
Memory size	17.5 KiB

Classification 1/3

- Probably easier
- Slightly imbalanced
  - Precision-Recall (instead of ROC Curve)

- 3 Models
  - XGBoost
  - SVC
  - LogReg

- 5-Fold Stratified CV

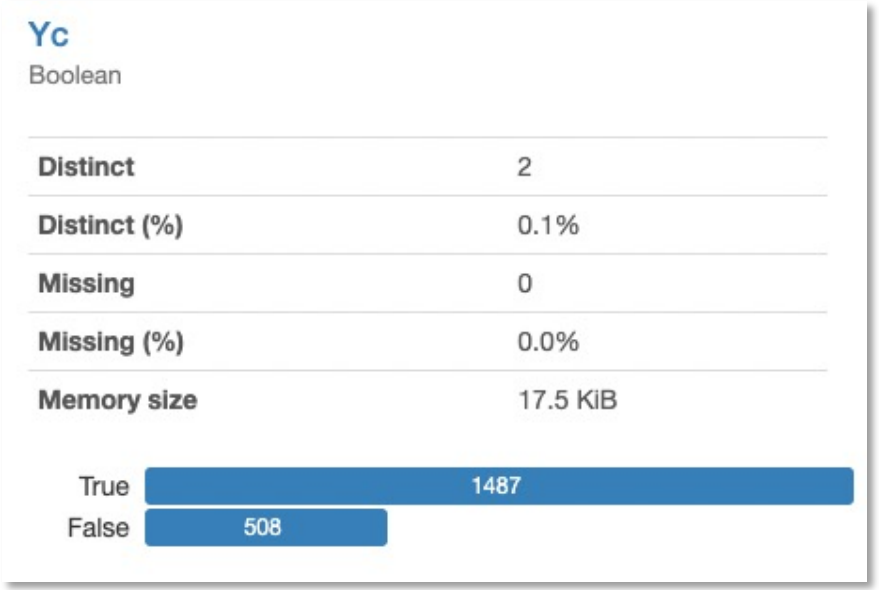

Classification

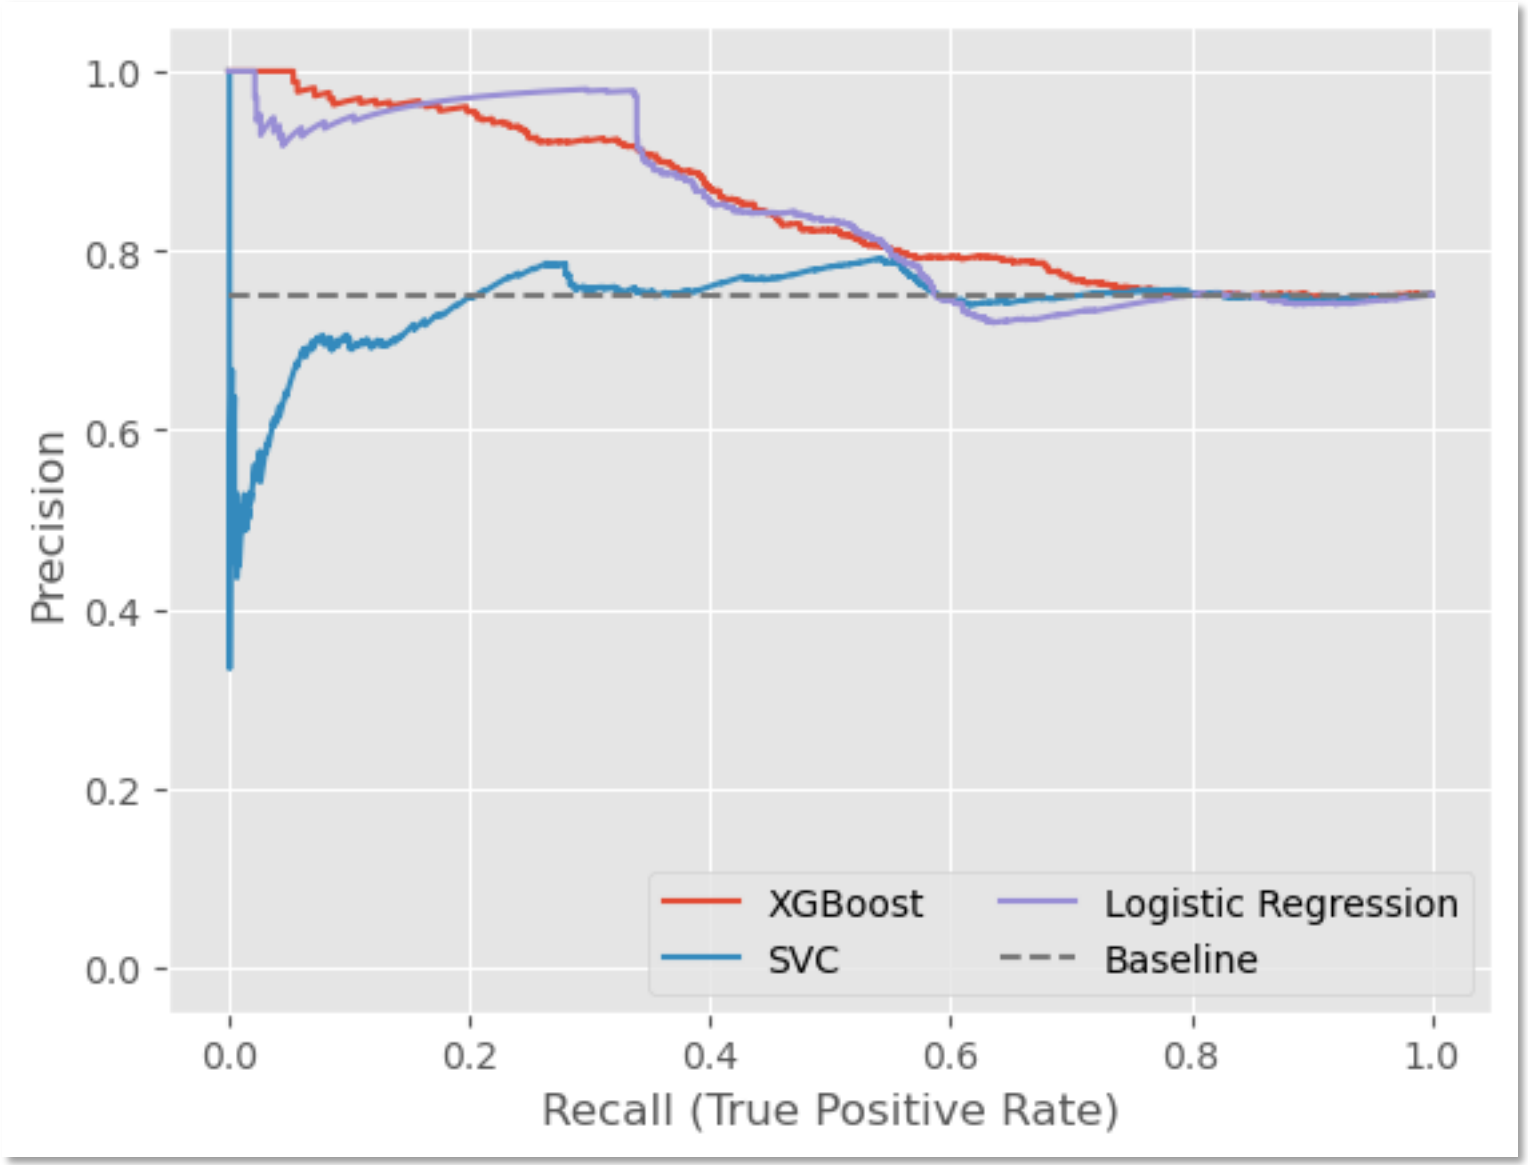

Classification

Conclusion 1/2

- Classification
  - LogReg or XGBoost do best (except at low thresholds)
  - XGBoost is the only one that stays above the baseline (“coin-toss”)
  - SVC isn’t doing very well
- Regression
  - Is really hard (target distribution is step-functiony), cf. horizontal stripes
- Overall, I’d go with XGBoost
  - It’s also “explainable”, so yay

Conclusion 2/2

- X4 seems redundant
  - Models don’t really change with/without

- Dropping one column from the one-hot encoded doesn’t matter
  - Models don’t change much with/without

What Now?

- Multi-Class Prediction?
- More Preprocessing
  - PCA? TDA?
  - Deal with noisy features? (X4, X5, X6)
    - → Fit Gaussian and take delta to be left with something interesting
    - → Or something more sophisticated (whitening / decorrelation)
    - Try dropping X6?
  - Drop one of one-hot? (Correlated)
- Pipeline in Sklearn
- Explainability